

Virtual Memory Report

Camille Williford - awillif4

November 8, 2021

The point of this lab was to familiarize students with how the xv6 simulator implements memory management. I implemented two system calls and prevented dereferencing null pointers for this lab. I am going to take a chronological order to this report detailing changes that I made to the xv6 OS in the order that I made them. Here is a link to the change log in GitHub: [Virtual Memory Commit](#)

1 Dereferencing Null Pointer

1.1 Makefile

Modified: 3 Lines

I changed the number of CPU's to 1, and made the entry point for page tables $0x1000$ per the lab write-up.

1.2 exec.c

Modified: 1 Lines

I changed the *exec()* to not create a page at address 0, but the given *PGSIZE*. This way user processes are not clobbering reserved memory.

1.3 vm.c

Modified: 1 Lines

I changed the *copyvm()* to not start copying memory at address 0, but the given *PGSIZE* i.e. the next page. I changed this function because it is responsible for copying a parent process's address state to a child. Now when *fork()* is called, the child process will have the correct address space.

2 System Calls

2.1 syscall.c

Modified: 4 Lines

Defined two new system calls *sys_mprotect()* and *sys_munprotect()* and their respective system functions.

2.2 syscall.h

Modified: 2 Lines

Assigned numbers to the new system calls *sys_mprotect()* and *sys_munprotect()*. These numbers are used when a user initiates a system call to safely determine which system call the user requested be performed.

2.3 user.h

Modified: 2 Lines

I prototyped the two new system calls as callable functions. These are the user functions that can be used to call into the operating system.

```
int mprotect(void *addr, int len);  
int munprotect(void *addr, int len);
```

2.4 usys.S

Modified: 2 Lines

Assembly code that loads register with syscall number, makes syscall, and makes transition to the kernel mode using int instruction. I added the two new system calls to this list.

2.5 sysproc.c

Modified: 28 Lines

I created the two new system calls *sys_mprotect()* and *sys_munprotect()*! These are a pass-through to the implementation in *vm.c*. Each parses the user input address with *argptr()* and length with *argint()*, and returns *-1* if unable to parse arguments.

2.6 vm.c

Modified: 2 Lines

Here the *mprotect()* and *munprotect()* system calls are implemented by securely changing the write bit on page table entries. These methods differ by only one line whether we are clearing the write bit in *mprotect()* or setting the write bit in *munprotect()*. Otherwise, I flush the TLB, check the user given length, and ensure the given address is page aligned. After the user input has been validated, I loop through all the pages within the given length and edit the write bit according to which function is used. Before editing the write bit, I ensure that page table entry is valid and that the user has permission to access that page table entry. If I cannot change the write bit, I return *-1*. Otherwise, I return 0 after looping through all the pages within a given length.

3 Tests

3.1 nullPtrTest.c

New File

This test simply creates a pointer to address 0x0, and then tries to get the value at that pointer.

```
#include "types.h"  
#include "stat.h"  
#include "user.h"  
  
int main(int argc, char **argv) {  
    uint* nullptr = (uint*) 0;  
    printf(1, "%x %x\n", nullptr, *nullptr);  
}
```

```
$  
$ nullPtrTest  
pid 11 nullPtrTest: trap 14 err 4 on cpu 0 eip 0x1004 addr 0x0--kill proc  
$
```

3.2 protectTest.c

New File

This test is more complex, testing that *mprotect()* and *munprotect()* work appropriately, so this test can be broken into two sections one for each system call. First I get the current process id, and then I initialize an address to protect and unprotect. I check *munprotect()* first by unprotecting a protected space and changing the value, which should not cause a trap. Second, I check that *mprotect()* properly protects an address space, by attempting to edit a protected value. This will cause a trap, since the protected space is uneditable.

```
int
main(int argc, char *argv[])
{
    int parent = getpid();
    // Create an address to protect
    int *val = (int *) sbrk(0);
    // Make sure the process memory size is at least a page
    sbrk(PGSIZE);
    // Protect the address
    mprotect(val, 1);

    if (fork() == 0) {
        munprotect(val, 1);
        *val = 5;
        printf(1, "TEST PASSED: unprotected value did not cause trap\n");
        exit();
    } else {
        wait();
    }

    if(fork() == 0) {
        sleep(5);
        *val = 5;
        printf(1, "TEST FAILED: protected value did not cause trap\n");
        kill(parent);
        exit();
    }

    // Parent: wait for child
    } else {
        wait();
    }
    printf(1, "TEST PASSED: protected value caused trap\n");
    exit();
}

$ protectTest
TEST PASSED: unprotected value did not cause trap
pid 18 protectTest: trap 14 err 7 on cpu 0 eip 0x1087 addr 0x4000--kill proc
TEST PASSED: protected value caused trap
$
```